

Roll Your Own SIEM

**ELK, Python, and Pattern Recognition
For Fun and Profit**

TASK June 2015

whoami

- my first PC was a VIC 20 (nobody could ever need more than 5k amiright?)



- current status: racking servers and scripting things at an awesome local ISP
- this talk is based around a project I did for my Ryerson compsec course, this field is new to me, so jump in – having a discussion around ML would be great since I'm primarily here to learn from you

So, This Is Happening



ATTACK ORIGINS

#	COUNTRY
6246	China
1436	United States
2306	Germany
2032	Russia
1689	Mil/Gov
1240	Ukraine
1886	South Korea
1132	Netherlands
867	Turkey
729	Japan

ATTACK TARGETS

#	COUNTRY
52381	United States
1242	Russia
208	Mil/Gov
131	Taiwan
93	Netherlands
25	Poland
22	France
21	Germany
3	United Kingdom
2	Italy

LIVE ATTACKS

ATTACKER			TARGET		TYPE
TIMESTAMP	ORGANIZATION	LOCATION	IP	LOCATION	PORT
2015-01-30 19:34:21.22	Hutchison Global	Central District, Hong	118.142.99.110	Saint Louis, United	smtp 25
2015-01-30 19:34:21.55	Chinanet Jiangsu Province	Nanjing, China	222.186.31.208	Seattle, United States	socks 1080
2015-01-30 19:34:21.93	ATAKOY-VAE-Pool	Istanbul, Turkey	46.1.145.211	Clifton, United States	unknown 60839
2015-01-30 19:34:22.23	YITAFENG	Zhaoyang, China	121.40.121.97	Seattle, United States	http 80
2015-01-30 19:34:22.58	NEC BIGLOBE Ltd.	Funabashi, Japan	122.135.34.205	Saint Louis, United	unknown 9791
2015-01-30 19:34:22.90	Webhosting.Net	Miami, United States	173.230.225.162	Miami, United States	unknown 21025
2015-01-30 19:34:23.28	Hong Kong Broadband	Central District, Hong	61.92.196.175	Saint Louis, United	unknown 49770
2015-01-30 19:34:23.65	Chinanet Jiangsu Province	Changzhou, China	61.160.224.129	Seattle, United States	rsync 873

ATTACK TYPES

#	SERVICE	PORT
8364	telnet	23
2789	sip	5060
2070	mysql	3306
1329	ms-wbt-server	3389
1808	http	80
1757	ms-sql-s	1433
1712	ssh	22
1432	active-net	3323



Blue Team – Traditional IDS

Signature based IDS has served for some time, but suffers from many problems – including collision attacks

“

And while most have already moved away from MD5, there is still a notable group that heavily uses this obsolete algorithm: **security vendors**. It seems that MD5 became the de-facto standard of fingerprinting malware samples and the industry doesn't seem to be willing to move away from this practice. Our friend Zoltán Balázs collected **a surprisingly long list of security vendors using MD5**, including the biggest names of the field.

The list includes for example Kaspersky, the discoverer of Flame who just recently reminded us that **MD5 is dead**, but just a few weeks earlier released a **report** including MD5 fingerprints only – ironically even the malware they analysed uses SHA-1 internally...

“Poisonous MD5 – Wolves Among the Sheep”
blog.silentsignal.eu

Insight?

Latest events

Level: 5 - Recipient domain is not found (450: Requested mail action not taken). 2015 Feb 02 11:47:46
Rule Id: 3303
Location: (syslog01) [REDACTED] /mnt/logs/smtp-outbound02/mail.log.2015.02.02
Src IP: [REDACTED]

Level: 5 - Recipient domain is not found (450: Requested mail action not taken). 2015 Feb 02 11:47:46
Rule Id: 3303
Location: (syslog01) [REDACTED] /mnt/logs/smtp-outbound02/mail.log.2015.02.02
Src IP: [REDACTED]

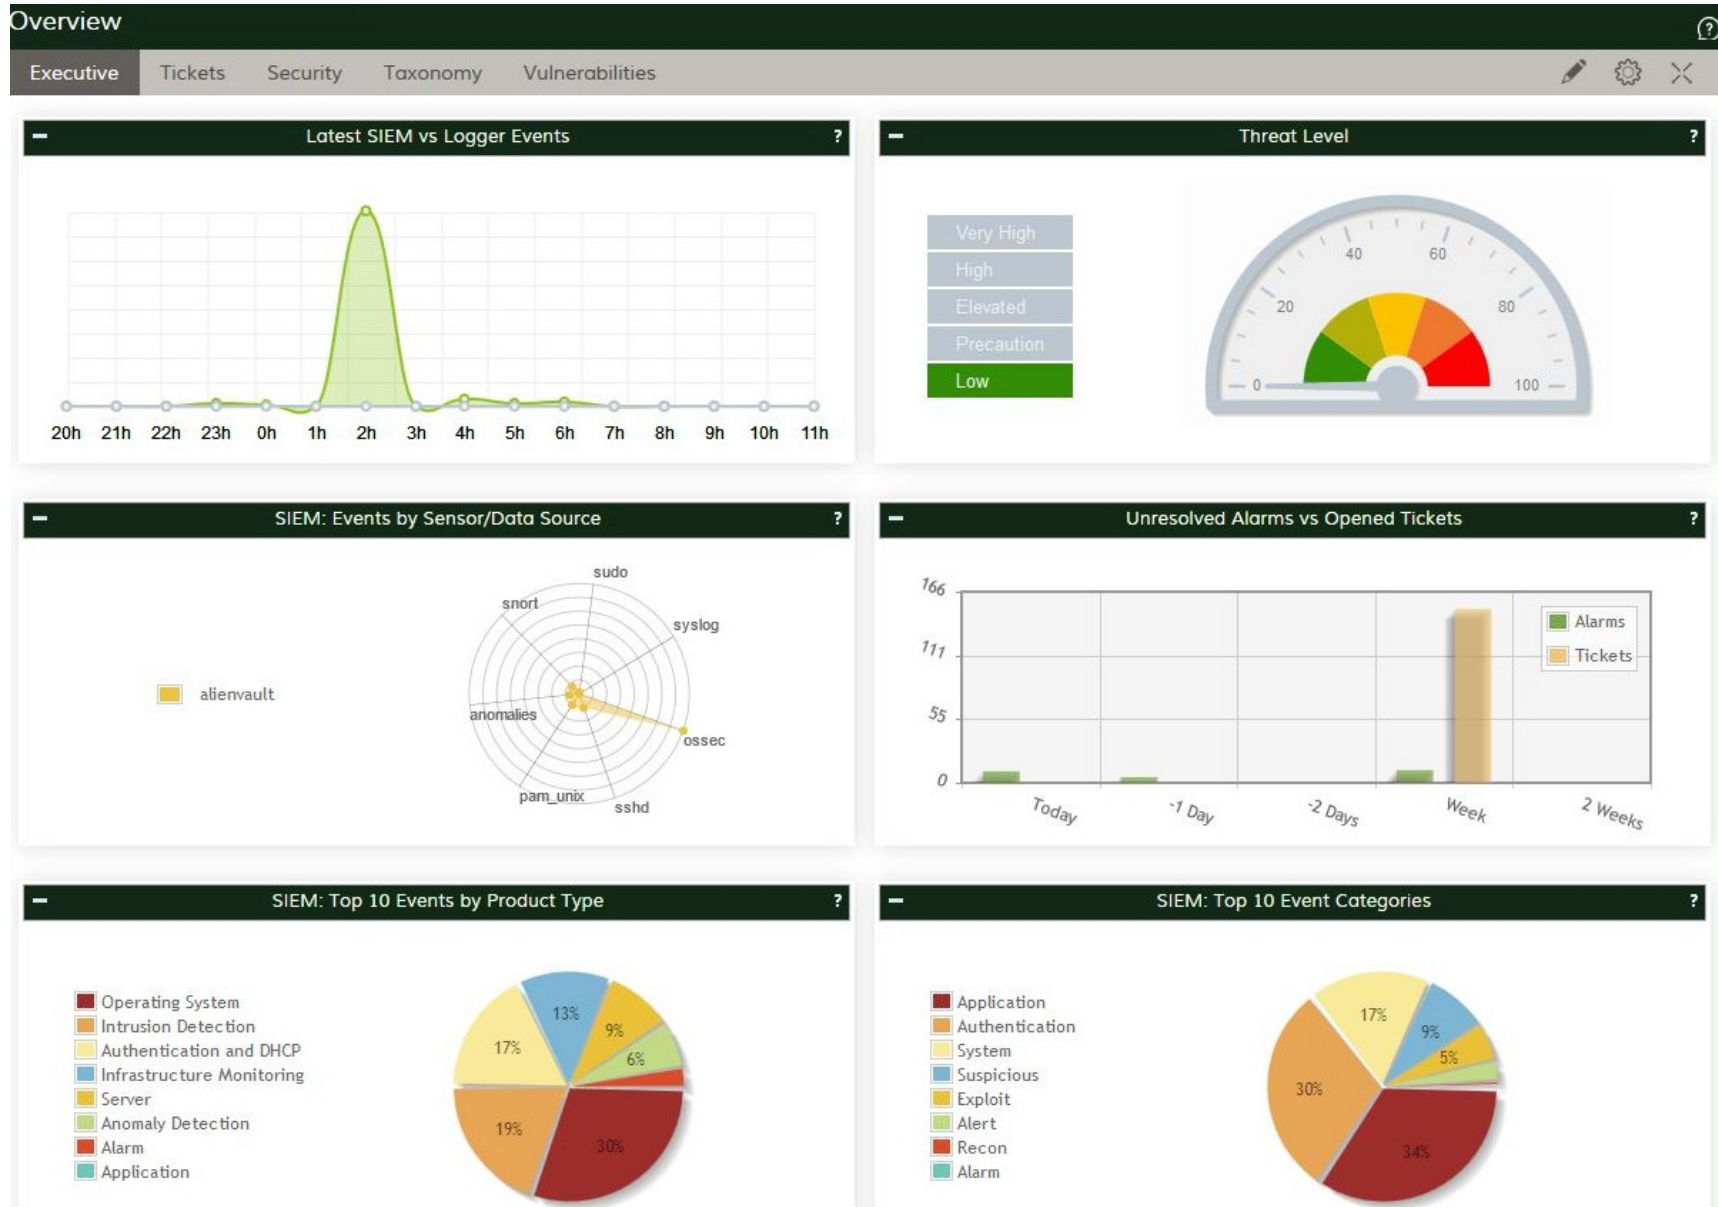
Level: 5 - Recipient domain is not found (450: Requested mail action not taken). 2015 Feb 02 11:47:46
Rule Id: 3303
Location: (syslog01) [REDACTED] /mnt/logs/smtp-outbound02/mail.log.2015.02.02
Src IP: [REDACTED]

Level: 5 - Recipient domain is not found (450: Requested mail action not taken). 2015 Feb 02 11:47:46
Rule Id: 3303
Location: (syslog01) [REDACTED] /mnt/logs/smtp-outbound02/mail.log.2015.02.02
Src IP: [REDACTED]

Level: 2 - Unknown problem somewhere in the system. 2015 Feb 02 11:47:46
Rule Id: 1002
Location: (syslog01) [REDACTED] /mnt/logs/smtp-auth05/mail.log.2015.02.02

Level: 5 - Recipient domain is not found (450: Requested mail action not taken). 2015 Feb 02 11:47:42
Rule Id: 3303
Location: (syslog01) [REDACTED] /mnt/logs/smtp-outbound02/mail.log.2015.02.02
Src IP: [REDACTED]

Insight?



Let's Make Our Own

But designing SIEM is not trivial, there are two problems to solve here

- Getting relevant intrusion information
 - Analysts overwhelmed by IDS alerting, false positives
- Presenting it in a useful way
 - Stock visualizations are full of “**chart junk**”, seldom match the way people process information or the context of your network
 - Always start with a question: What problem are you solving?

Dashboards: First Principles

Information Design – Tufte, Few – The practice of presenting information in a way that fosters efficient and effective understanding of it – Wikipedia

“...working memory is limited to three or four simultaneous chunks of information at a time”

– Stephen Few

*(Why Do We Visualize Quantitative Data?,
<http://www.perceptualedge.com/blog/?p=1897>)*

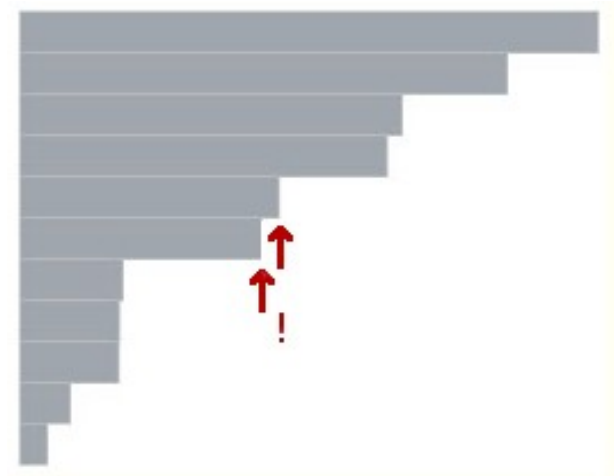
Design: Pie Charts Are Evil

Things people are bad at (not exhaustive)

- Estimating 2D areas



<- compare ->



- Holding lists in working memory

Less is More: The Crow Epistemology

Quick, how many dots?



How about here?



ELK Stack – Open Source and Flexible



[downloads](#) [docs](#) [support](#) [discuss](#) [contact](#)

[PRODUCTS](#)

[SUBSCRIPTIONS](#)

[LEARN](#)

[COMMUNITY](#)

[USE CASES](#)

[BLOG](#)

[ABOUT](#)



[GO](#)

Kibana | Explore & Visualize Your Data




See the Value in Your Data

- Flexible analytics and visualization platform
- Real-time summary and charting of streaming data
- Intuitive interface for a variety of users
- Instant sharing and embedding of dashboards

[DOWNLOAD](#)

[SUBSCRIPTIONS](#)

Get Product Updates



Seamless Integration with Elasticsearch

Architected to work with Elasticsearch, Kibana gives shape to any kind of data — structured and unstructured — indexed into Elasticsearch. It also benefits from Elasticsearch's powerful search and analytics capabilities.



Give Shape to Your Data

To better understand large volumes of data, easily create bar charts, line and scatter plots, histograms, pie charts, and maps.

KIBANA

Visualization front
end (JS)

ELASTICSEARCH

LOGSTASH

Parses, stores
logs, runs on JVM

REDIS

Key-value cache
for scalability

SHIPPER(S)

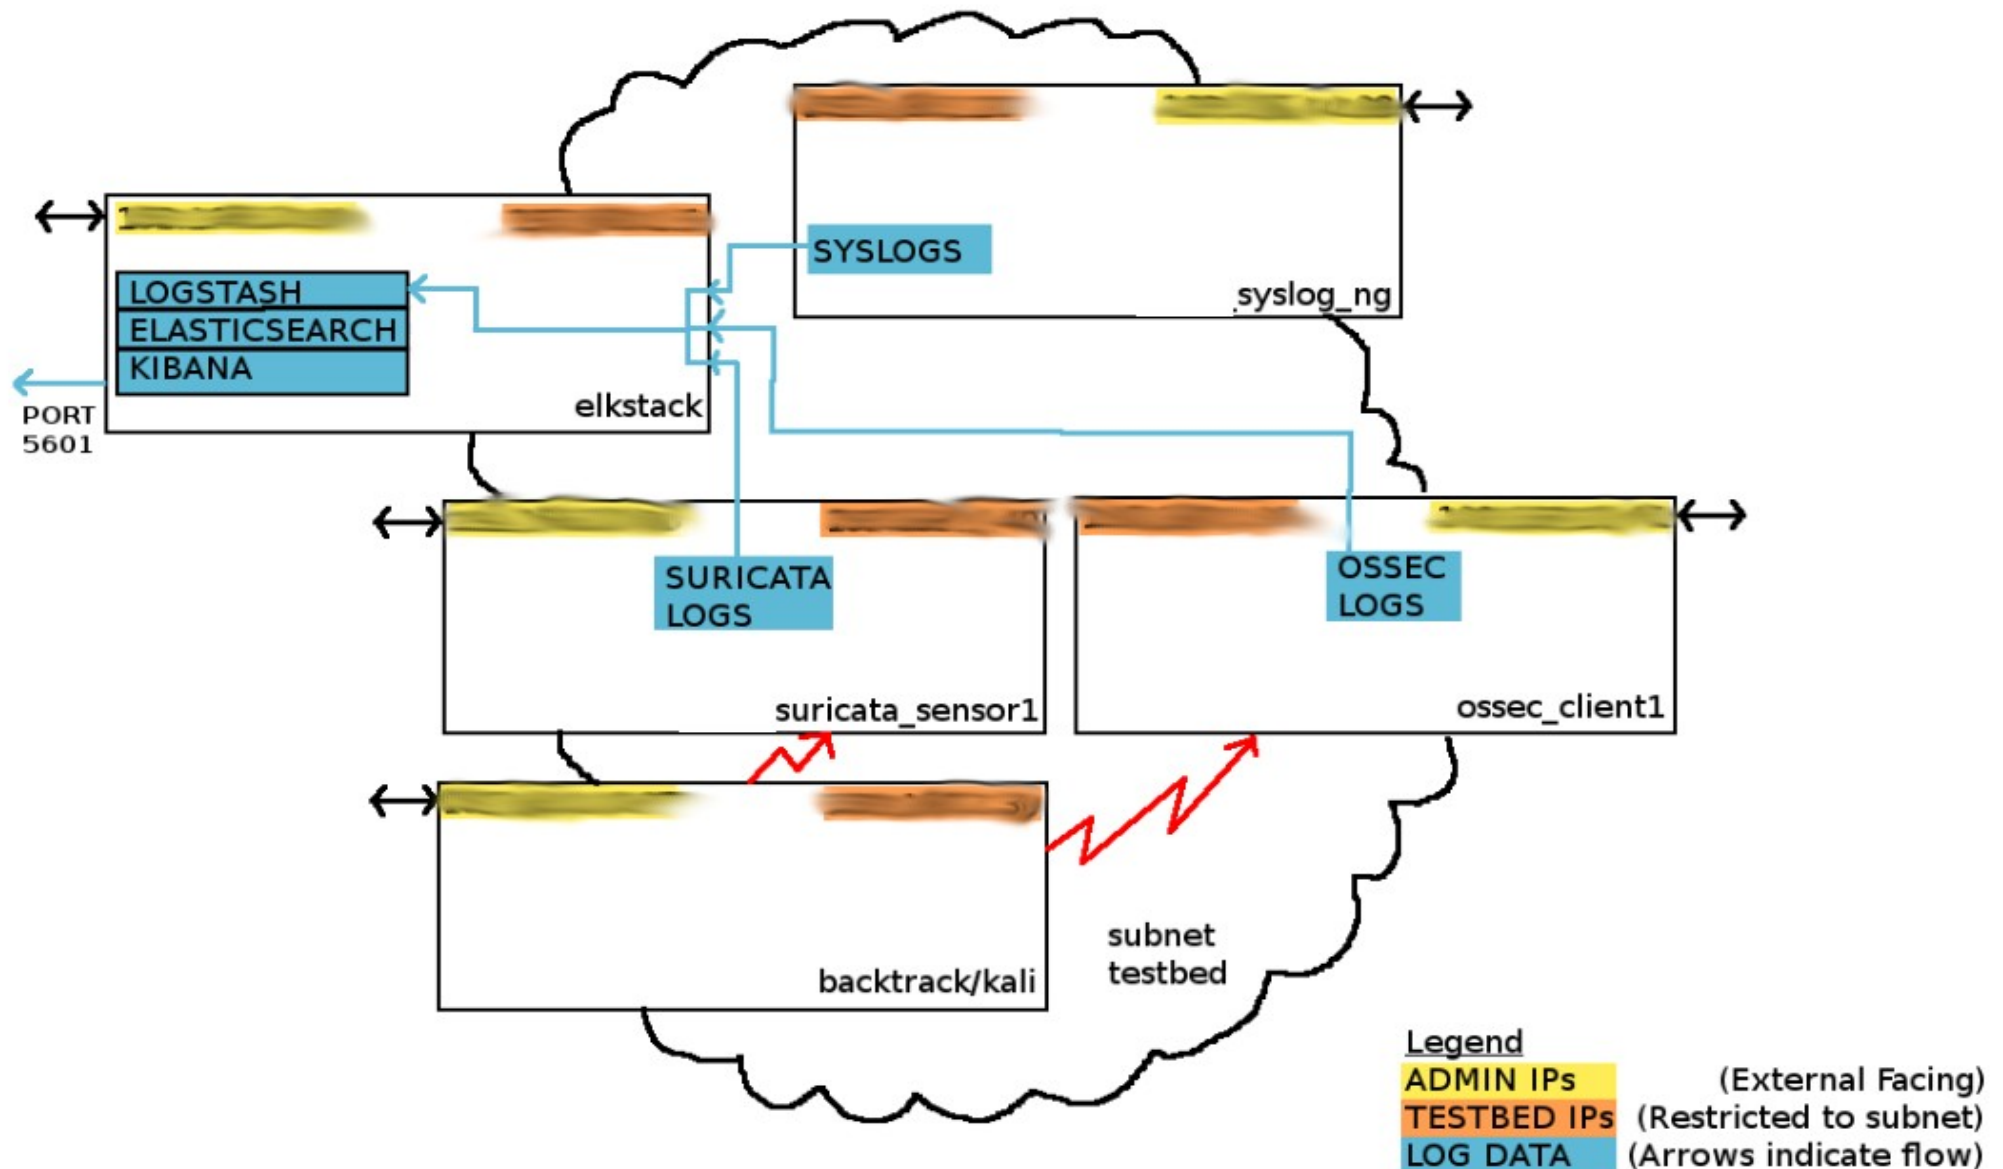


COLLECT



ALL THE THINGS!

Sample Testbed




A close-up, top-down view of a large pile of cut logs. The logs are of various diameters and are stacked closely together, creating a textured, circular pattern. The wood is a warm, golden-brown color, with some logs showing distinct growth rings and others having darker, possibly charred or weathered, surfaces. The lighting is even, highlighting the natural grain and textures of the wood.

NOW WHAT?

PATTERN RECOGNITION




SCIENCE



One Approach - Fast Fourier Transform


```
19  
20 import numpy  
21 from scipy import fftpack  
22 import matplotlib.pyplot as plot  
23
```



Time Series Data as a Waveform

Syslog Events - Severity Type 5

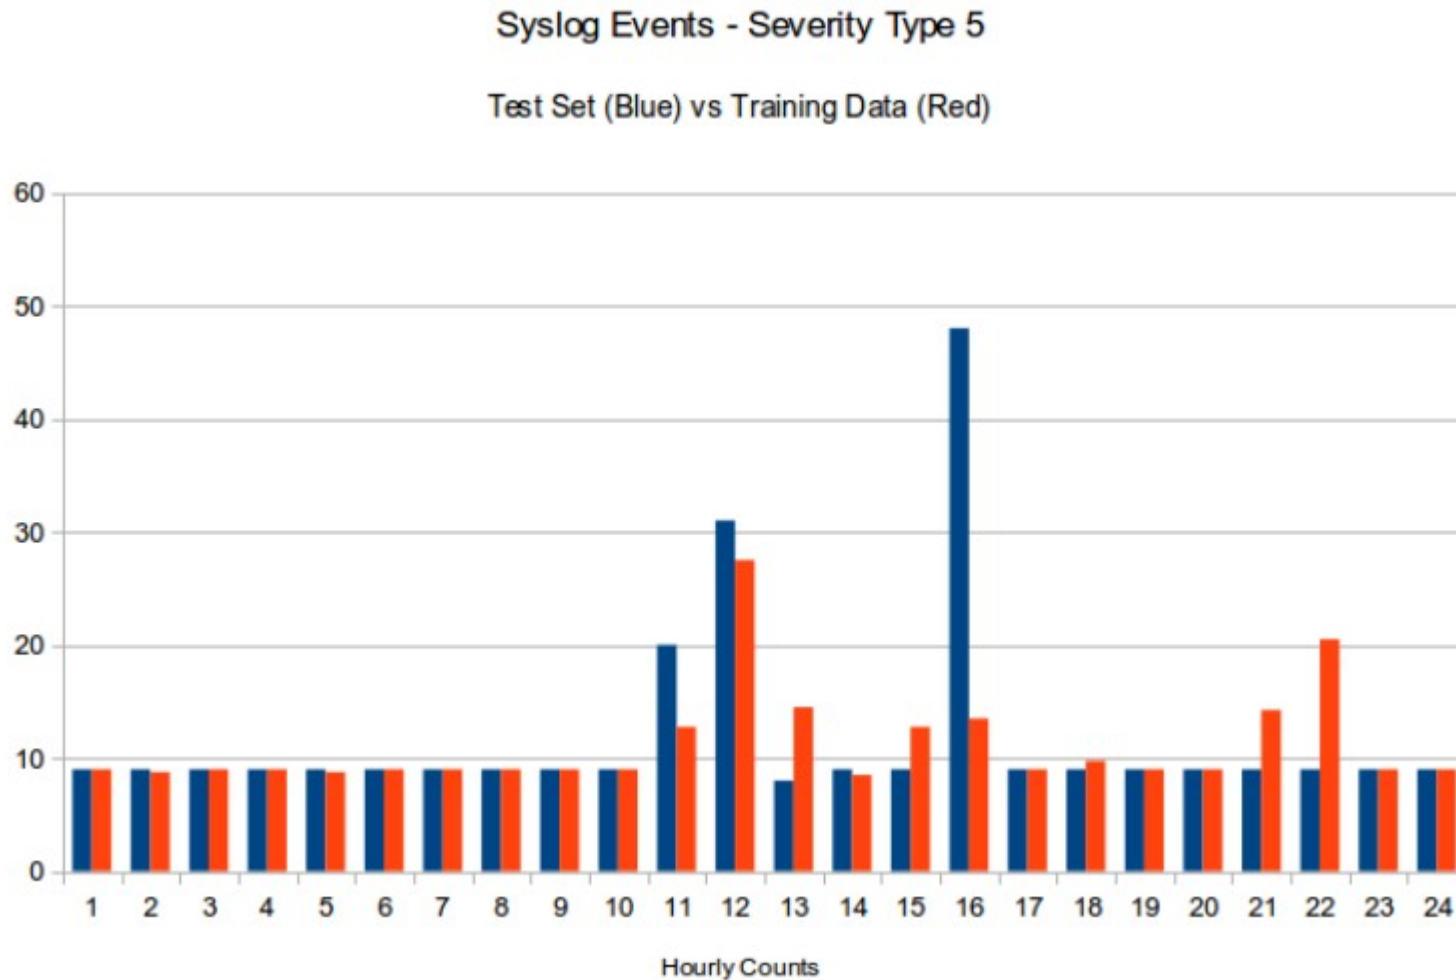
Testbed Results for April 1, 2015



Anomaly Detection

- Any time series (logs) may be viewed as a waveform
- Complex waveforms may be decomposed into simpler components, fingerprinted
- It's not hard to use off the shelf machine learning libraries to experiment with creating fingerprints of your own network traffic
- Machine learning in its simplest form is the comparison of test data to a set of training data
- In this case, I used a smoothed 30-day history of traffic by day and alert type to create a training dataset
- Any significant differences get flagged, tagged and reinserted into logstash

Test Set Data vs Training Set Data



Smooth Data

ALERT on test data:

```
[ 9 9 9 9 9 9 9 9 9 9 9 20 31 8 9 9 48 9 9 9 9 9
9 9 9]
```

vs training data

```
[ 9. 8.75 9. 9. 8.75 9. 9. 9. 9. 9. 12.75 27.5 14.5 8.5 12.75
13.5 9. 9.75 9. 9. 14.25 20.5 9. 9. ]
```

For logtype syslog, subtype 5 for date 2015-04-01, the value 20 at hour 10 may be an outlier!

Test/Train difference: 0.538976326886

Observation variance: 0.252585149809

ALERT on test data:

```
[ 9 9 9 9 9 9 9 9 9 9 9 20 31 8 9 9 48 9 9 9 9 9
9 9 9]
```

vs training data

```
[ 9. 8.75 9. 9. 8.75 9. 9. 9. 9. 9. 12.75 27.5 14.5 8.5 12.75
13.5 9. 9.75 9. 9. 14.25 20.5 9. 9. ]
```

For logtype syslog, subtype 5 for date 2015-04-01, the value 48 at hour 15 may be an outlier!

Test/Train difference: 0.538976326886

Observation variance: 0.44736895924

OUTLIER RESULTS found for 2015-04-01 by type, subtype in the following hours:

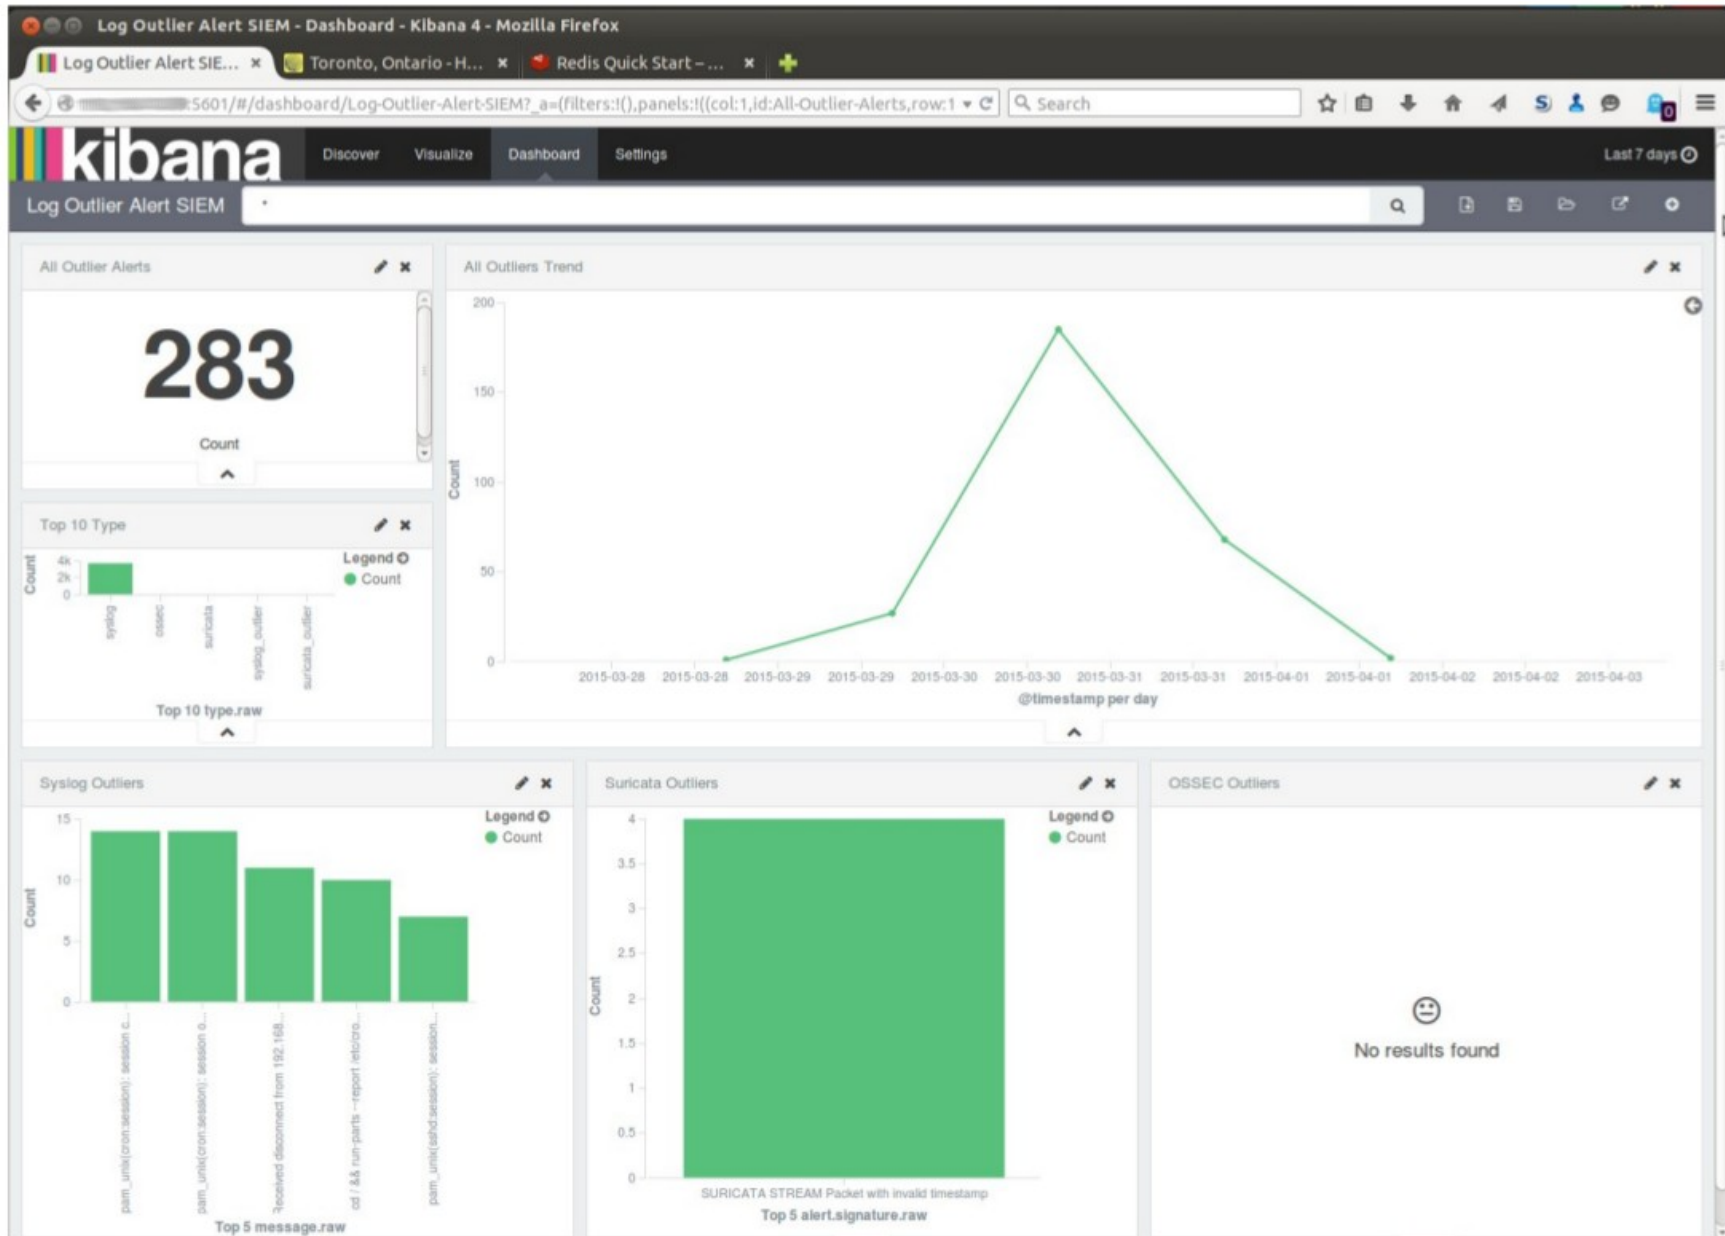
```
{'syslog': {'5': [10, 15]}}
```

Results of reindexing outlying records

A total of 68 records were updated and saved as file

'records2015-04-01'

Display



Demo(?)

Let us propitiate the demo gods
May they be merciful



esquery.py

```
$ ./esquery.py
usage: one of the arguments -a/--action -q/--query is required
usage: esquery.py [-h]
                  (-a [{stringquery,termquery,termsagg,count}] | -q QUERY)
                  (-l | -t TERMS) [-d [HOST]] [-i [INDEX]] [-s [SIZE]]
                  [-r RANGE RANGE RANGE] [-f FIELDS [FIELDS ...]] [-c]

esquery 0.1

optional arguments:
  -h, --help                show this help message and exit
  -a [{stringquery,termquery,termsagg,count}], --action [{stringquery,termquery,termsagg,count}]
                           actions wrap common query types (default:stringquery)
                           - mutually exclusive of --query
  -q QUERY, --query QUERY   raw elasticsearch json query
  -l, --list                display available indices - mutually exclusive of
                           --terms
  -t TERMS, --terms TERMS   some text to query - format for stringquery is STRING,
                           format for termquery is TERM:STRING, format for
                           termsagg is TERM
  -d [HOST], --host [HOST]  the elasticsearch host IP (default: localhost)
  -i [INDEX], --index [INDEX]
                           specifies a specific index to query (default:all)
  -s [SIZE], --size [SIZE]  number of hits to return (default: 10)
  -r RANGE RANGE RANGE, --range RANGE RANGE RANGE
                           range filter, specify field then beginning and end
                           points as numeric arguments or in YYYY-MM-DD format
                           for dates eg; --range severity 5 9 or -r timestamp
                           2015-03-05 2015-03-11
  -f FIELDS [FIELDS ...], --fields FIELDS [FIELDS ...]
                           specify source fields to include in search
                           (default:all fields)
  -c, --count               return a count of hits only
```


Lets Look at the Logs We've Collected

```
$ ./esquery.py -l -a
Namespace(action=None, count=False, fields=None, host='127.0.0.1', index='logstash*', list=True, query=None, range=None, size=None, terms=None)
health status index pri rep docs.count docs.deleted store.size pri.store.size
yellow open logstash-2015.06.14 5 1 251 120 246.4kb 246.4kb
yellow open logstash-2015.05.27 5 1 1880 0 839.9kb 839.9kb
yellow open logstash-2015.04.06 5 1 245 0 317.8kb 317.8kb
yellow open logstash-2015.05.04 5 1 1846 0 875.1kb 875.1kb
yellow open logstash-2015.04.05 5 1 248 0 252.7kb 252.7kb
yellow open logstash-2015.04.20 5 1 14600 0 1.5mb 1.5mb
yellow open logstash-2015.03.09 5 1 248 0 379kb 379kb
yellow open logstash-2015.05.08 5 1 182 0 227.7kb 227.7kb
yellow open logstash-2015.04.14 5 1 14610 0 1.6mb 1.6mb
yellow open logstash-2015.03.23 5 1 280 31 410.9kb 410.9kb
yellow open .kibana 1 1 38 1 65.3kb 65.3kb
yellow open logstash-2015.06.22 5 1 281 0 233.3kb 233.3kb
yellow open logstash-2015.04.22 5 1 8445 0 1mb 1mb
yellow open logstash-2015.03.16 5 1 261 0 304.5kb 304.5kb
yellow open logstash-2015.03.18 5 1 245 0 324.9kb 324.9kb
yellow open logstash-2015.04.11 5 1 245 0 186.6kb 186.6kb
yellow open logstash-2015.06.18 5 1 243 0 189.7kb 189.7kb
yellow open logstash-2015.04.21 5 1 14601 0 1.6mb 1.6mb
yellow open logstash-2015.02.16 5 1 7298966 0 3gb 3gb
yellow open logstash-2015.04.18 5 1 14598 0 1.5mb 1.5mb
yellow open logstash-2015.03.14 5 1 268 0 273.5kb 273.5kb
yellow open logstash-2015.03.05 5 1 24963 0 4.4mb 4.4mb
yellow open logstash-2015.04.07 5 1 256 0 239.6kb 239.6kb
yellow open logstash-2015.03.19 5 1 247 0 306.7kb 306.7kb
yellow open logstash-2015.03.07 5 1 244 0 270.5kb 270.5kb
yellow open logstash-2015.03.11 5 1 280 0 311.3kb 311.3kb
yellow open logstash-2015.06.15 5 1 259 0 154.1kb 154.1kb
yellow open logstash-2015.06.21 5 1 274 0 251.5kb 251.5kb
yellow open logstash-2015.03.27 5 1 538 193 511.2kb 511.2kb
yellow open logstash-2015.04.19 5 1 14604 0 1.6mb 1.6mb
```

Query a Single Syslog?

```
./esquery.py --host ... --action termquery --terms type:syslog --size 1
Namespace(action='termquery', count=False, fields=None, host='...', index='logstash*', list=False, query=None, range=None, size='1', terms=['type:syslog'])
http://.../logstash*/_search -d {"size": "1", "query": {"term": {"type": "syslog"}}}
{
  "_shards": {
    "failed": 0,
    "successful": 370,
    "total": 370
  },
  "hits": {
    "hits": [
      {
        "_id": "AUvrJ7PCAztdK20Uh9n6",
        "_index": "logstash-2015.03.05",
        "_score": 3.7593648,
        "_source": {
          "@timestamp": "2015-03-05T18:17:01.000Z",
          "@version": "1",
          "host": "...",
          "logsource": "...",
          "message": "pam_unix(cron:session): session opened for user root by (uid=0)",
          "pam_by": "(uid=0)",
          "pam_caller": "cron:session",
          "pam_module": "pam_unix",
          "pam_session_state": "opened",
          "path": "/var/log/auth.log",
          "pid": "4015",
          "program": "CRON",
          "syslog_facility": "user-level",
          "syslog_facility_code": 1,
          "syslog_severity": "notice",
          "syslog_severity_code": 5,
          "tags": [
            "_grokparsefailure"
          ],
          "type": "syslog",
          "username": "root"
        }
      }
    ]
  }
}
```

Start Kibana

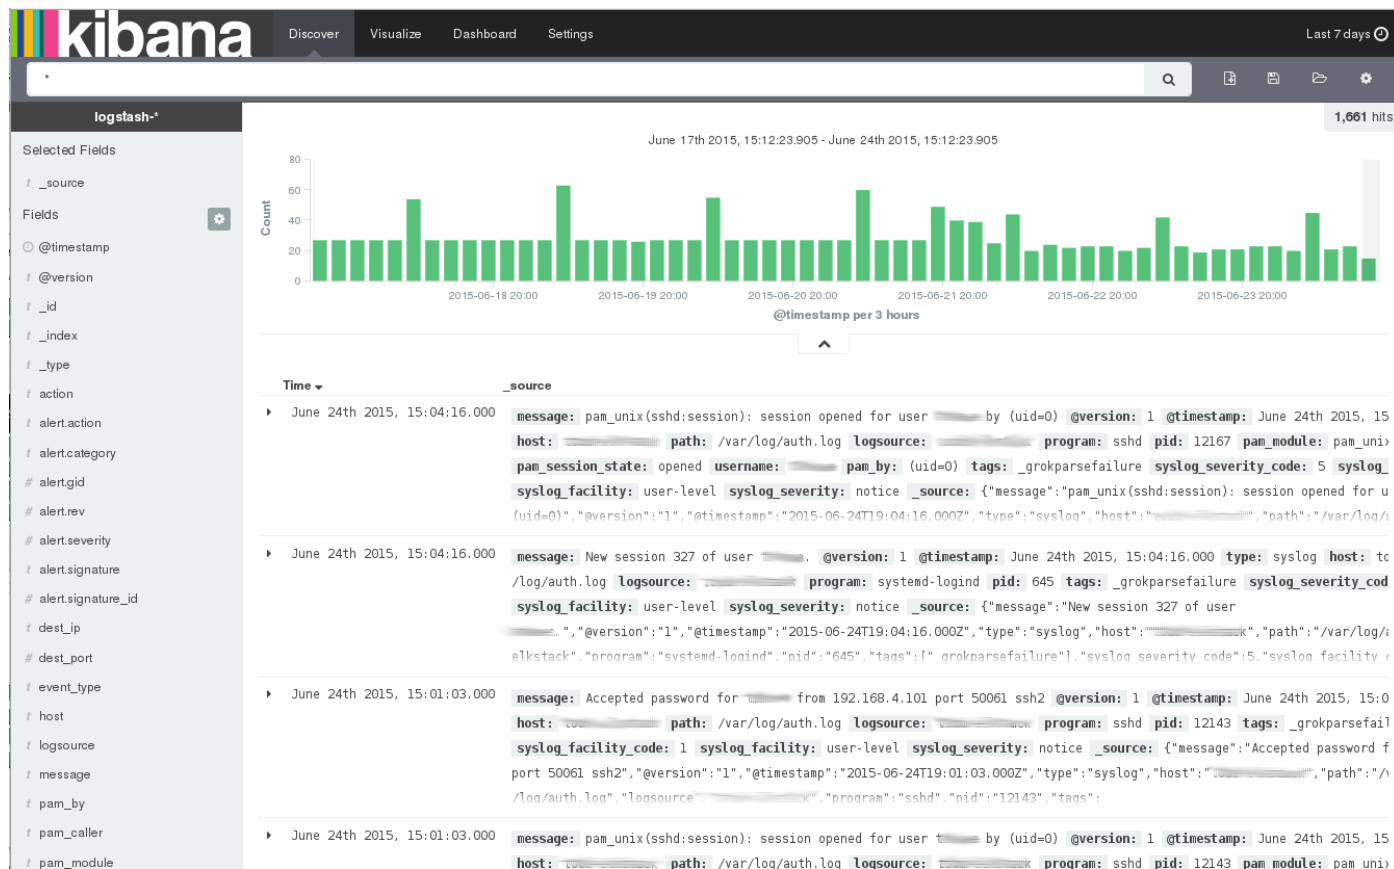
```
~$ cd Downloads/
~/Downloads$ ls
elasticsearch-1.4.4.deb      postfix-grok-patterns-master
kibana-4.0.0-linux-x64      postfix-grok-patterns-master.zip
kibana-4.0.0-linux-x64.tar.gz  test-postfix-logs
logstash_1.4.2-1-2c0f5a1_all.deb

~/Downloads$ cd kibana-4.0.0-linux-x64/
~/Downloads/kibana-4.0.0-linux-x64$ ls
bin  config  LICENSE.txt  node  plugins  README.txt  src

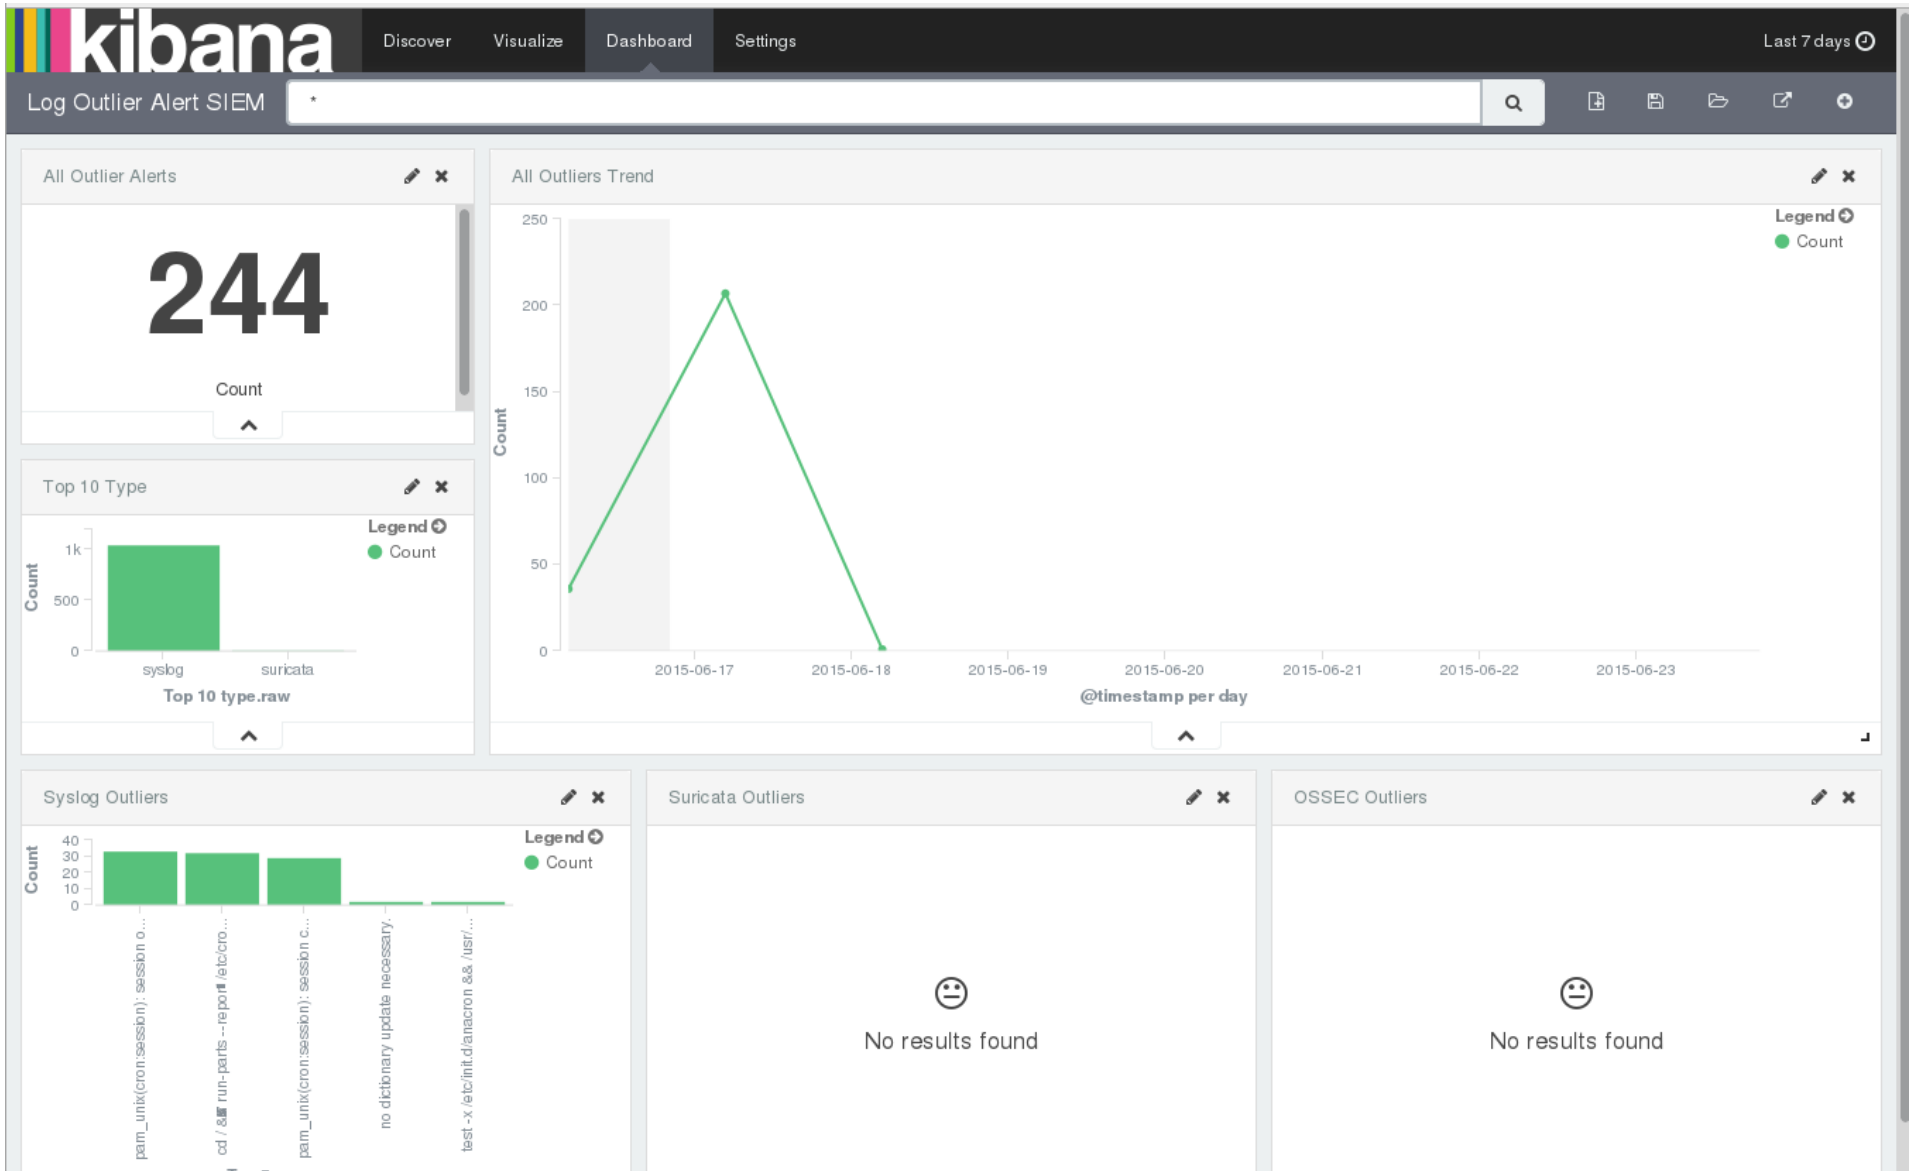
~/Downloads/kibana-4.0.0-linux-x64$ cd bin
~/Downloads/kibana-4.0.0-linux-x64/bin$ ./kibana
{"@timestamp":"2015-06-24T19:08:57.614Z","level":"info","message":"Listening on 0.0.0.0:5601","node_env":"production"}
```

Launch the Kibana GUI

- Connect on port 5601
- Kibana spews a bunch of JSON to the console for debugging, and the browser displays the Discover view



Dashboard View – Pre Analysis Jobs



Normalize, Vectorize, Analyze and Store

- set TEST_DAY in esconstants.py to the number of days ago the day you want to analyze was
- script currently only operates on an entire day
- then run ./escontrol.py
- uncaught exceptions (oops) mean sharding failed temporarily on Elasticsearch, try again or add a new node to the cluster

Busy Busy JVM

tehowe@todd-elkstack (192.168.10.89) - byobu

File Edit View Search Terminal Help

```
1  [|||||] 4.3% Tasks: 71, 176 thr; 1 running
2  [|||||] 28.8% Load average: 0.69 0.60 0.61
3  [|||||] 10.9% Uptime: 12 days, 05:29:05
4  [|||||] 4.3%
Mem[|||||] 1768/3965MB
Swp[|||||] 0/1292MB
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
705	elasticse	20	0	6731M	2172M	1022M	S	48.8	54.8	29h31:01	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1018	elasticse	20	0	6731M	2172M	1022M	S	31.3	54.8	16h11:44	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1008	elasticse	20	0	6731M	2172M	1022M	S	3.8	54.8	1h43:06	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1011	elasticse	20	0	6731M	2172M	1022M	S	3.3	54.8	1h43:06	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1013	elasticse	20	0	6731M	2172M	1022M	S	3.3	54.8	1h43:06	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1015	elasticse	20	0	6731M	2172M	1022M	S	3.3	54.8	1h43:06	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1023	elasticse	20	0	6731M	2172M	1022M	S	3.3	54.8	2h50:01	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
17383		20	0	25012	4216	2920	R	0.9	0.1	0:00.45	htop
12212		20	0	35116	12284	2324	S	0.5	0.3	0:02.58	SCREEN -S byobu -c /usr/share/byobu/profiles/byoburc /usr/bin/byobu-shell
1033	logstash	39	19	2697M	210M	19692	S	0.5	5.3	10:43.88	/usr/bin/java -Djava.io.tmpdir=/var/lib/logstash -Xmx500m -XX:+UseParNewGC -XX:+UseConcMarkSweepGC -Djava.awt.
1142	elasticse	20	0	6731M	2172M	1022M	S	0.5	54.8	1:26.81	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1100	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	51:18.73	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
669	logstash	39	19	2697M	210M	19692	S	0.0	5.3	26:21.22	/usr/bin/java -Djava.io.tmpdir=/var/lib/logstash -Xmx500m -XX:+UseParNewGC -XX:+UseConcMarkSweepGC -Djava.awt.
958	www-data	20	0	91532	3592	2120	S	0.0	0.1	1:21.17	nginx: worker process
671	redis	20	0	38532	3176	2360	S	0.0	0.1	12:18.94	/usr/bin/redis-server 192.168.1.5:6379
1039	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	10:39.39	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
11611	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	0:01.15	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1226	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	21:04.82	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
9771	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	7:16.07	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1227	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	21:05.25	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1225	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	21:04.41	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1152	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	1:29.60	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1094	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	2:31.02	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1187	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	1:30.48	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1193	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	1:29.68	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
9685		20	0	99484	3980	3044	S	0.0	0.1	0:00.19	sshd: pts/1
1167	logstash	39	19	2697M	210M	19692	S	0.0	5.3	4:10.37	/usr/bin/java -Djava.io.tmpdir=/var/lib/logstash -Xmx500m -XX:+UseParNewGC -XX:+UseConcMarkSweepGC -Djava.awt.
1022	logstash	39	19	2697M	210M	19692	S	0.0	5.3	0:30.26	/usr/bin/java -Djava.io.tmpdir=/var/lib/logstash -Xmx500m -XX:+UseParNewGC -XX:+UseConcMarkSweepGC -Djava.awt.
1141	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	1:28.48	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1210	logstash	39	19	2697M	210M	19692	S	0.0	5.3	2:18.52	/usr/bin/java -Djava.io.tmpdir=/var/lib/logstash -Xmx500m -XX:+UseParNewGC -XX:+UseConcMarkSweepGC -Djava.awt.
1132	logstash	39	19	2697M	210M	19692	S	0.0	5.3	3:22.25	/usr/bin/java -Djava.io.tmpdir=/var/lib/logstash -Xmx500m -XX:+UseParNewGC -XX:+UseConcMarkSweepGC -Djava.awt.
1192	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	1:29.62	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1149	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	1:27.79	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1121	Debian-gd	20	0	1468M	129M	66608	S	0.0	3.3	0:28.49	gnome-shell --mode=gdm
1144	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	1:30.46	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC
1145	elasticse	20	0	6731M	2172M	1022M	S	0.0	54.8	1:29.76	/usr/lib/jvm/java-7-openjdk-amd64/bin/java -Xms256m -Xmx1g -Xss256k -Djava.awt.headless=true -XX:+UseParNewGC

F1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice F8Nice F9Kill F10Quit

0\$ bash 1-\$ kibana 2*\$ top

@ Debian GNU/Linux 8.0

Menu: <F9> 2# 12d5h 0.69 4x2.36GHz 3.9GB45% 2015-06-24 16:21:48

Results!

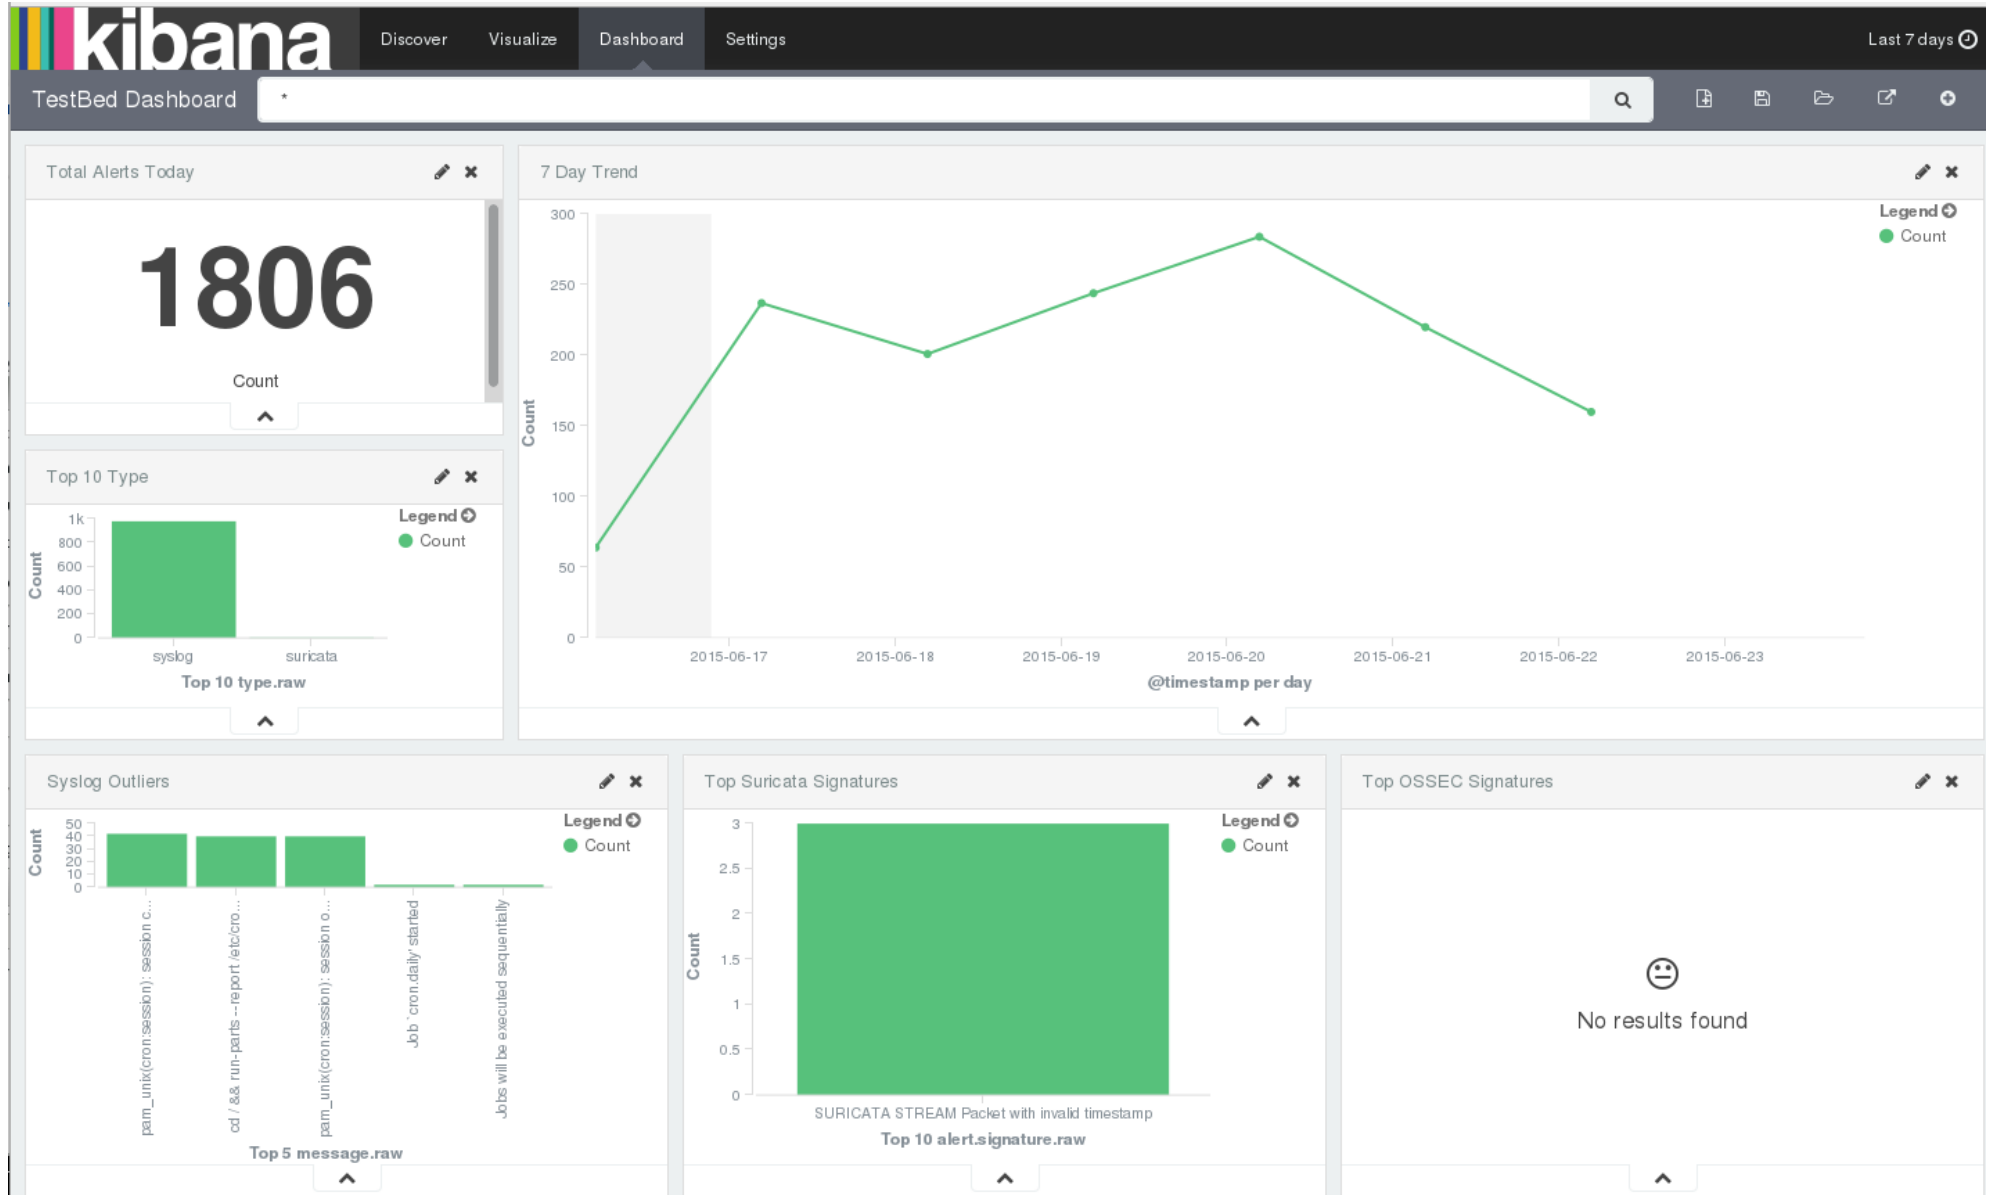
```
+-----$ ./escontrol.py
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-19", "gte": "2015-06-19"}}}}}, "size": 1000000}
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-18", "gte": "2015-06-18"}}}}}, "size": 10000}
Got day: 2015-06-18
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-17", "gte": "2015-06-17"}}}}}, "size": 10000}
Got day: 2015-06-17
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-16", "gte": "2015-06-16"}}}}}, "size": 10000}
Got day: 2015-06-16
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-15", "gte": "2015-06-15"}}}}}, "size": 10000}
Got day: 2015-06-15
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-14", "gte": "2015-06-14"}}}}}, "size": 10000}
Got day: 2015-06-14
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-13", "gte": "2015-06-13"}}}}}, "size": 10000}
Got day: 2015-06-13
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-12", "gte": "2015-06-12"}}}}}, "size": 10000}
Got day: 2015-06-12
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-11", "gte": "2015-06-11"}}}}}, "size": 10000}
Dropped day: 2015-06-11
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-10", "gte": "2015-06-10"}}}}}, "size": 10000}
Dropped day: 2015-06-10
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-09", "gte": "2015-06-09"}}}}}, "size": 10000}
Dropped day: 2015-06-09
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-08", "gte": "2015-06-08"}}}}}, "size": 10000}
Dropped day: 2015-06-08
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-06-07", "gte": "2015-06-07"}}}}}, "size": 10000}
Dropped day: 2015-06-07
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-05-20", "gte": "2015-05-20"}}}}}, "size": 10000}
Dropped day: 2015-05-20
http://...:9200/logstash/_search -d {"query": {"filtered": {"query": {"query_string": {"query": "*"}}, "filter": {"range": {"@timestamp": {"lte": "2015-05-19", "gte": "2015-05-19"}}}}}, "size": 10000}
Dropped day: 2015-05-19

WARNING Subtype 2210044 doesn't exist in training data for 2210044
*** Sending this as an alert to SIEM

OUTLIER RESULTS found for 2015-06-19 by type, subtype in the following hours:
{'suricata': {'2210044': [6, 21]}}

Results of reindexing outlying records
-----
A total of 1 records were updated and saved as file 'records2015-06-19'
```


Dashboard View – Post Analysis Jobs





Please Enjoy Responsibly


https://github.com/z3r0fox/ML_log_analysis



toy ML log analysis with SciPy and Elasticsearch — Edit

 **2** commits



 **1** branch








 **0** releases


 **1** contributor

 branch: **master** ▾ **ML_log_analysis** / + 

POC log analysis suite

 **z3r0fox** authored 2 minutes ago latest commit 26f87e75b0 

 README.md	POC log analysis suite	2 minutes ago
 esanalyze.py	POC log analysis suite	2 minutes ago
 esconstants.py	POC log analysis suite	2 minutes ago
 escontrol.py	POC log analysis suite	2 minutes ago
 esindex.py	POC log analysis suite	2 minutes ago
 esquery.py	POC log analysis suite	2 minutes ago
 esvectorize.py	POC log analysis suite	2 minutes ago

 **README.md**

Interesting Reads

- Wikipedia: Information Design
https://en.wikipedia.org/wiki/Information_design
- Stephen Few, Why Do We Visualize Quantitative Data?
<http://www.perceptualedge.com/blog/?p=1897>
- Managing Logstash with the Redis Client
<http://www.nightbluefruit.com/blog/2014/03/managing-logstash-with-the-redis-client/>
- An interactive introduction to FFT
<http://betterexplained.com/articles/an-interactive-guide-to-the-fourier-transform/>
- Data Driven Security (Jay Jacobs & Bob Rudis)
BONUS: Pan-fried gnocchi
<http://datadrivensecurity.info/>

Thank You



@z3r0fox